

Assignment No-9

Name : Shubham Khalate

Class :S.Y-A

Roll No.: 23107062

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv("TCS.csv")

df
   Close \ Date      Open      High      Low
0    2004-08-27  122.800003  122.800003  119.820000  120.332497
1    2004-08-30  121.237503  123.750000  120.625000  123.345001
2    2004-08-31  123.312500  123.750000  122.000000  123.512497
3    2004-09-01  123.750000  124.375000  122.949997  123.487503
4    2004-09-02  123.737503  125.574997  123.250000  124.207497
...
4489  2022-10-18  3150.000000  3155.350098  3128.550049  3144.699951
4490  2022-10-19  3159.000000  3159.000000  3112.000000  3121.850098
4491  2022-10-20  3105.000000  3160.000000  3105.000000  3157.300049
4492  2022-10-21  3157.800049  3160.399902  3127.000000  3137.399902
4493  2022-10-24  3170.100098  3178.000000  3155.000000  3161.699951

   Adj Close      Volume
0     88.088272  30646000.0
1     90.293549  24465208.0
2     90.416122  21194656.0
3     90.397820  19935544.0
```

```

4      90.924896  21356352.0
...
4489  3144.699951  1793722.0
4490  3121.850098  1194289.0
4491  3157.300049  1587601.0
4492  3137.399902  1021913.0
4493  3161.699951  260949.0

[4494 rows x 7 columns]

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4494 entries, 0 to 4493
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Date        4494 non-null    object 
 1   Open         4486 non-null    float64
 2   High         4486 non-null    float64
 3   Low          4486 non-null    float64
 4   Close        4486 non-null    float64
 5   Adj Close   4486 non-null    float64
 6   Volume       4486 non-null    float64
dtypes: float64(6), object(1)
memory usage: 245.9+ KB

df.describe()

      Open           High           Low          Close      Adj Close
\count  4486.000000  4486.000000  4486.000000  4486.000000  4486.000000
mean    1146.182768  1158.538059  1132.825428  1145.521462  1049.456965
std     994.070086  1003.010607  984.043404  993.346465  992.062558
min     112.000000  116.112503  103.837502  111.550003  86.565590
25%    290.693748  295.300010  285.931259  290.275009  219.902748
50%    977.450012  995.000000  970.250000  981.337524  820.627839
75%    1564.774963  1597.287476  1548.662537  1576.781219  1443.631561
max    4033.949951  4043.000000  3980.000000  4019.149902  3964.502686

      Volume
\count  4.486000e+03
mean    3.620596e+06

```

```
std      3.162368e+06
min      0.000000e+00
25%     1.942489e+06
50%     2.804749e+06
75%     4.297410e+06
max     8.806715e+07

df.ndim
2

df.dtypes
Date        object
Open       float64
High       float64
Low        float64
Close      float64
Adj Close   float64
Volume     float64
dtype: object

df.size
31458

df.shape
(4494, 7)

df.head()

```

	Date	Open	High	Low	Close	Adj
0	2004-08-27	122.800003	122.800003	119.820000	120.332497	88.088272
1	2004-08-30	121.237503	123.750000	120.625000	123.345001	90.293549
2	2004-08-31	123.312500	123.750000	122.000000	123.512497	90.416122
3	2004-09-01	123.750000	124.375000	122.949997	123.487503	90.397820
4	2004-09-02	123.737503	125.574997	123.250000	124.207497	90.924896

	Volume
0	30646000.0
1	24465208.0
2	21194656.0
3	19935544.0
4	21356352.0

```
df.tail()
```

	Date	Open	High	Low	
Close \					
4489	2022-10-18	3150.000000	3155.350098	3128.550049	3144.699951
4490	2022-10-19	3159.000000	3159.000000	3112.000000	3121.850098
4491	2022-10-20	3105.000000	3160.000000	3105.000000	3157.300049
4492	2022-10-21	3157.800049	3160.399902	3127.000000	3137.399902
4493	2022-10-24	3170.100098	3178.000000	3155.000000	3161.699951
	Adj Close	Volume			
4489	3144.699951	1793722.0			
4490	3121.850098	1194289.0			
4491	3157.300049	1587601.0			
4492	3137.399902	1021913.0			
4493	3161.699951	260949.0			

```
df.isnull()
```

	Date	Open	High	Low	Close	Adj Close	Volume	
0	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False
...
4489	False	False	False	False	False	False	False	False
4490	False	False	False	False	False	False	False	False
4491	False	False	False	False	False	False	False	False
4492	False	False	False	False	False	False	False	False
4493	False	False	False	False	False	False	False	False

[4494 rows x 7 columns]

```
df.isnull().value_counts()
```

Date	Open	High	Low	Close	Adj Close	Volume	
False	False	False	False	False	False	False	4486
	True	True	True	True	True	True	8

Name: count, dtype: int64

```
df.isnull().sum()
```

Date	0
Open	8
High	8

```
Low      8  
Close     8  
Adj Close 8  
Volume    8  
dtype: int64
```

```
df.notnull()
```

```
Date  Open  High  Low   Close  Adj Close  Volume  
0    True  True  True  True   True    True    True  
1    True  True  True  True   True    True    True  
2    True  True  True  True   True    True    True  
3    True  True  True  True   True    True    True  
4    True  True  True  True   True    True    True  
...  ...  ...  ...  ...  ...  ...  ...  
4489  True  True  True  True   True    True    True  
4490  True  True  True  True   True    True    True  
4491  True  True  True  True   True    True    True  
4492  True  True  True  True   True    True    True  
4493  True  True  True  True   True    True    True
```

```
[4494 rows x 7 columns]
```

```
df.notnull().value_counts()
```

```
Date  Open  High  Low   Close  Adj Close  Volume  
True  True  True  True   True    True    True    4486  
      False False False False  False   False  False     8  
Name: count, dtype: int64
```

```
df.notnull().sum()
```

```
Date      4494  
Open      4486  
High      4486  
Low       4486  
Close     4486  
Adj Close 4486  
Volume    4486  
dtype: int64
```

```
df.dropna(inplace=True)
```

```
df.isnull().sum()
```

```
Date      0  
Open      0  
High      0  
Low       0  
Close     0  
Adj Close 0
```

```

Volume      0
dtype: int64

df[ "Date" ] = pd.to_datetime(df[ "Date" ])

df.info()

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 4486 entries, 2004-08-27 to 2022-10-24
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Open        4486 non-null    float64 
 1   High       4486 non-null    float64 
 2   Low         4486 non-null    float64 
 3   Close       4486 non-null    float64 
 4   Adj Close   4486 non-null    float64 
 5   Volume      4486 non-null    float64 
dtypes: float64(6)
memory usage: 245.3 KB

df.set_index('Date', inplace=True)

df.head()

          Open        High        Low     Close  Adj Close
\Date
2004-08-27  122.800003  122.800003  119.820000  120.332497  88.088272
2004-08-30  121.237503  123.750000  120.625000  123.345001  90.293549
2004-08-31  123.312500  123.750000  122.000000  123.512497  90.416122
2004-09-01  123.750000  124.375000  122.949997  123.487503  90.397820
2004-09-02  123.737503  125.574997  123.250000  124.207497  90.924896

          Volume
Date
2004-08-27  30646000.0
2004-08-30  24465208.0
2004-08-31  21194656.0
2004-09-01  19935544.0
2004-09-02  21356352.0

plt.figure(figsize=(12, 5))
plt.plot(df.index, df[ "Close" ], label='Close Price')
plt.title('Close Price Over Time')

```

```

plt.xlabel('Date')
plt.ylabel('Close Price')
plt.grid()
plt.show()

```



```

from statsmodels.tsa.stattools import adfuller

def adf_test(series):
    result = adfuller(series)
    labels = ["ADF Statistics:", "P-Value", "Lags_used", "Number Of Observation"]
    for value,label in zip(result,labels):
        print(label, ' : ', str(value))
        if(result[1]<= 0.05):
            print("Strong Evidences against the Null hypothesis,
reject the null hypothesis, and it stationary also has no unit root")
        else:
            print("Weak evidences against the null hypothesis ,Time series has a root, indicates it is not stationary")

adf_test(df['Close'])

ADF Statistics: : 0.8059405233011762
Weak evidences against the null hypothesis ,Time series has a root,
indicates it is not stationary
P-Value : 0.9917314565589413
Weak evidences against the null hypothesis ,Time series has a root,
indicates it is not stationary
Lags_used : 32
Weak evidences against the null hypothesis ,Time series has a root,
indicates it is not stationary
Number Of Observation : 4453

```

```
Weak evidences against the null hypothesis ,Time series has a root,  
indicates it is not stationary
```

```
df['Diff of Close'] = df['Close']-df['Close'].shift(1)  
df.head()
```

	Open	High	Low	Close	Adj Close
Date					
2004-08-27	122.800003	122.800003	119.820000	120.332497	88.088272
2004-08-30	121.237503	123.750000	120.625000	123.345001	90.293549
2004-08-31	123.312500	123.750000	122.000000	123.512497	90.416122
2004-09-01	123.750000	124.375000	122.949997	123.487503	90.397820
2004-09-02	123.737503	125.574997	123.250000	124.207497	90.924896

	Volume	Diff of Close	12 Diff
Date			
2004-08-27	30646000.0	NaN	NaN
2004-08-30	24465208.0	3.012504	NaN
2004-08-31	21194656.0	0.167496	NaN
2004-09-01	19935544.0	-0.024994	NaN
2004-09-02	21356352.0	0.719994	NaN

```
df['12 Diff'] = df['Close']-df['Close'].shift(12)
```

```
adf_test(df['12 Diff'].dropna())
```

ADF Statistics: : -12.17533024797796

Strong Evidences against the Null hypothesis, reject the null hypothesis, and it stationary also has no unit root

P-Value : 1.3904621898083073e-22

Strong Evidences against the Null hypothesis, reject the null hypothesis, and it stationary also has no unit root

Lags_used : 31

Strong Evidences against the Null hypothesis, reject the null hypothesis, and it stationary also has no unit root

Number Of Observation : 4442

Strong Evidences against the Null hypothesis, reject the null hypothesis, and it stationary also has no unit root

```
adf_test(df['Diff of Close'].dropna())
```

ADF Statistics: : -14.990218644441317

Strong Evidences against the Null hypothesis, reject the null hypothesis, and it stationary also has no unit root

```

P-Value : 1.1282227743588002e-27
Strong Evidences against the Null hypothesis, reject the null
hypothesis, and it stationary also has no unit root
Lags_used : 31
Strong Evidences against the Null hypothesis, reject the null
hypothesis, and it stationary also has no unit root
Number Of Observation : 4453
Strong Evidences against the Null hypothesis, reject the null
hypothesis, and it stationary also has no unit root

plt.figure(figsize=(14, 7))
plt.plot(df.index, df['Diff of Close'], label='Differenced Close
Price',color='orange')
plt.title('Differenced Close Price Over Time')
plt.xlabel('Date')
plt.ylabel('Differenced Close Price')
plt.legend()
plt.show()


```




```

from statsmodels.tsa.seasonal import seasonal_decompose
decompose = seasonal_decompose(df['Diff of
Close'].dropna(),model='addictive',period=30)
decompose.plot()


```





```
from statsmodels.graphics.tsaplots import plot_acf,plot_pacf
pd.plotting.autocorrelation_plot(df['Close'])
```

```
<Axes: xlabel='Lag', ylabel='Autocorrelation'>
```




```

from statsmodels.graphics.tsaplots import plot_acf,plot_pacf
from statsmodels.tsa.stattools import acf,pacf

acf1 = plot_acf(df['Diff of Close'].dropna(),alpha = 0.05,lags = 32)
"""
Correlation with each month and there is Pearson Correlation
From Acf we find the parameter q
Q < blue area(error)
"""

'\nCorrelation with each month and there is Pearson Correlation\nFrom
Acf we find the parameter q\nQ < blue area(error)\n'

```




```
xacf = pd.DataFrame(acf(df['Diff of Close'].dropna()))
print(xacf)
```

```
          0
0    1.000000
1   -0.016574
2   -0.036398
3   -0.016025
4    0.004304
5    0.048556
6   -0.010082
7   -0.026128
8   -0.021521
9    0.011150
10  -0.002779
11  -0.024594
12   0.016758
13   0.014371
14  -0.060926
15   0.017095
16   0.009561
17   0.011891
18  -0.034102
19  -0.003790
```

```
20  0.046849
21 -0.013866
22  0.029597
23 -0.028894
24 -0.007285
25 -0.011749
26 -0.051880
27  0.004981
28 -0.024287
29 -0.014058
30 -0.039046
31 -0.078959
32 -0.010349
33  0.012727
34 -0.005207
35  0.027834
36 -0.010172


acf12 = plot_acf(df['12_Diff'].dropna())
```




```
xacf = pd.DataFrame(acf(df['12_Diff'].dropna()))
print(xacf)
```

```
          0
0  1.000000
1  0.909188
2  0.820225
3  0.731854
4  0.649476
5  0.564754
6  0.470228
7  0.373339
8  0.285304
9  0.206131
10 0.122287
11 0.037955
12 -0.045308
13 -0.040620
14 -0.040647
15 -0.037193
16 -0.038258
17 -0.042232
18 -0.045696
19 -0.047362
20 -0.057994
21 -0.080545
22 -0.098546
23 -0.122595
24 -0.141596
25 -0.158140
26 -0.172177
27 -0.185589
28 -0.199821
29 -0.205880
30 -0.209829
31 -0.209777
32 -0.192381
33 -0.165422
34 -0.141144
35 -0.112478
36 -0.094657
```

```
pacf = plot_pacf(df['Diff of Close'].dropna())
```



```
pacf12= plot_pacf(df['12 Diff'].dropna())
```



df		Open	High	Low	Close	Adj
Close \ Date						
2004-08-27	122.800003	122.800003	119.820000	120.332497		
88.088272						
2004-08-30	121.237503	123.750000	120.625000	123.345001		
90.293549						
2004-08-31	123.312500	123.750000	122.000000	123.512497		
90.416122						
2004-09-01	123.750000	124.375000	122.949997	123.487503		
90.397820						
2004-09-02	123.737503	125.574997	123.250000	124.207497		
90.924896						
...
...						
2022-10-18	3150.000000	3155.350098	3128.550049	3144.699951		
3144.699951						
2022-10-19	3159.000000	3159.000000	3112.000000	3121.850098		
3121.850098						
2022-10-20	3105.000000	3160.000000	3105.000000	3157.300049		
3157.300049						

```
2022-10-21 3157.800049 3160.399902 3127.000000 3137.399902  
3137.399902  
2022-10-24 3170.100098 3178.000000 3155.000000 3161.699951  
3161.699951
```

```
          Volume Diff of Close    12 Diff  
Date  
2004-08-27 30646000.0      NaN      NaN  
2004-08-30 24465208.0      3.012504  NaN  
2004-08-31 21194656.0      0.167496  NaN  
2004-09-01 19935544.0      -0.024994  NaN  
2004-09-02 21356352.0      0.719994  NaN  
...       ...   ...     ...  
2022-10-18 1793722.0      32.949951  147.399902  
2022-10-19 1194289.0      -22.849853  117.300049  
2022-10-20 1587601.0      35.449951  172.350098  
2022-10-21 1021913.0      -19.900147  46.250000  
2022-10-24 260949.0       24.300049  59.750000
```

[4486 rows x 8 columns]


```
from statsmodels.tsa.arima.model import ARIMA  
train_size = int(len(df) * 0.8)  
train, test = df.iloc[:train_size], df.iloc[train_size:]  
  
model = ARIMA(train["Close"].dropna(), order=(1,1,1))  
model_fit = model.fit()  
  
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\  
tsa_model.py:473: ValueWarning: A date index has been provided, but it  
has no associated frequency information and so will be ignored when  
e.g. forecasting.  
    self._init_dates(dates, freq)  
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\  
tsa_model.py:473: ValueWarning: A date index has been provided, but it  
has no associated frequency information and so will be ignored when  
e.g. forecasting.  
    self._init_dates(dates, freq)  
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\  
tsa_model.py:473: ValueWarning: A date index has been provided, but it  
has no associated frequency information and so will be ignored when  
e.g. forecasting.  
    self._init_dates(dates, freq)  
  
forecast = model_fit.forecast(steps=len(test))  
plt.figure(figsize=(14,7))  
plt.plot(train.index, train["Close"], label='Train', color="#203147")  
plt.plot(test.index, test["Close"], label='Test', color="#01ef63")  
plt.plot(test.index, forecast, label='Forecast', color='orange')  
plt.title('Close Price Forecast')
```

```

plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()

C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:836: ValueWarning: No supported index is available.
Prediction results will be given with an integer index beginning at
`start`.
    return get_prediction_index(
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:836: FutureWarning: No supported index is available. In
the next version, calling this method in a model without a supported
index will result in an exception.
    return get_prediction_index()

```



```

print(f"AIC: {model_fit.aic}")
print(f"BIC: {model_fit.bic}")

AIC: 29303.565323136136
BIC: 29322.12053756873

print("Mean Close Price:", df['Close'].mean())
print("Max Close Price:", df['Close'].max())
print("Min Close Price:", df['Close'].min())

Mean Close Price: 1145.5214623165402
Max Close Price: 4019.149902
Min Close Price: 111.550003

```

```

from sklearn.metrics import mean_squared_error
import numpy as np

rmse = np.sqrt(mean_squared_error(test["Close"], forecast))
print("RMSE:", rmse)

RMSE: 1027.441327038619

```

SARIMA

```

from statsmodels.tsa.statespace.sarimax import SARIMAX
import matplotlib.pyplot as plt

model = SARIMAX(train["Close"].dropna(), order=(1,1,1),
                 seasonal_order=(1,1,1,12))
# 12 = yearly seasonality for monthly data
model_fit = model.fit()


C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:473: ValueWarning: A date index has been provided, but it
has no associated frequency information and so will be ignored when
e.g. forecasting.
    self._init_dates(dates, freq)
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:473: ValueWarning: A date index has been provided, but it
has no associated frequency information and so will be ignored when
e.g. forecasting.
    self._init_dates(dates, freq)

forecast = model_fit.forecast(steps=len(test))


C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:836: ValueWarning: No supported index is available.
Prediction results will be given with an integer index beginning at
`start`.
    return get_prediction_index()
C:\ProgramData\anaconda3\Lib\site-packages\statsmodels\tsa\base\
tsa_model.py:836: FutureWarning: No supported index is available. In
the next version, calling this method in a model without a supported
index will result in an exception.
    return get_prediction_index()


plt.figure(figsize=(14,7))
plt.plot(train.index, train["Close"], label='Train', color="#203147")
plt.plot(test.index, test["Close"], label='Test', color="#01ef63")
plt.plot(test.index, forecast, label='Forecast', color='orange')
plt.title('Close Price Forecast using SARIMA')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()

```



```
model_fit.plot_diagnostics(figsize=(15, 12))
```





```
from sklearn.metrics import mean_squared_error
import numpy as np

rmse = np.sqrt(mean_squared_error(test["Close"], forecast))
print("RMSE:", rmse)

RMSE: 749.4810675385451
```

Prophet

```
df = pd.read_csv('TCS.CSV')

df['Date'] = pd.to_datetime(df['Date'])
df = df[['Date', 'Close']]

!pip install Prophet

Collecting Prophet
  Downloading prophet-1.1.6-py3-none-win_amd64.whl.metadata (3.6 kB)
Collecting cmdstanpy>=1.0.4 (from Prophet)
```

```
  Downloading cmdstanpy-1.2.5-py3-none-any.whl.metadata (4.0 kB)
Requirement already satisfied: numpy>=1.15.4 in c:\programdata\anaconda3\lib\site-packages (from Prophet) (1.26.4)
Requirement already satisfied: matplotlib>=2.0.0 in c:\programdata\anaconda3\lib\site-packages (from Prophet) (3.9.2)
Requirement already satisfied: pandas>=1.0.4 in c:\programdata\anaconda3\lib\site-packages (from Prophet) (2.2.2)
Collecting holidays<1,>=0.25 (from Prophet)
  Downloading holidays-0.70-py3-none-any.whl.metadata (34 kB)
Requirement already satisfied: tqdm>=4.36.1 in c:\programdata\anaconda3\lib\site-packages (from Prophet) (4.66.5)
Collecting importlib-resources (from Prophet)
  Downloading importlib_resources-6.5.2-py3-none-any.whl.metadata (3.9 kB)
Collecting stadio<2.0.0,>=0.4.0 (from cmdstanpy>=1.0.4->Prophet)
  Downloading stadio-0.5.1-py3-none-any.whl.metadata (1.6 kB)
Requirement already satisfied: python-dateutil in c:\programdata\anaconda3\lib\site-packages (from holidays<1,>=0.25->Prophet) (2.9.0.post0)
Requirement already satisfied: contourpy>=1.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.0.0->Prophet) (1.2.0)
Requirement already satisfied: cycler>=0.10 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.0.0->Prophet) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.0.0->Prophet) (4.51.0)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.0.0->Prophet) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.0.0->Prophet) (24.1)
Requirement already satisfied: pillow>=8 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.0.0->Prophet) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib>=2.0.0->Prophet) (3.1.2)
Requirement already satisfied: pytz>=2020.1 in c:\programdata\anaconda3\lib\site-packages (from pandas>=1.0.4->Prophet) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in c:\programdata\anaconda3\lib\site-packages (from pandas>=1.0.4->Prophet) (2023.3)
Requirement already satisfied: colorama in c:\programdata\anaconda3\lib\site-packages (from tqdm>=4.36.1->Prophet) (0.4.6)
Requirement already satisfied: six>=1.5 in c:\programdata\anaconda3\lib\site-packages (from python-dateutil->holidays<1,>=0.25->Prophet) (1.16.0)
  Downloading prophet-1.1.6-py3-none-win_amd64.whl (13.3 MB)
    ----- 0.0/13.3 MB ? eta -:-:-
    ----- 0.5/13.3 MB 2.8 MB/s eta
0:00:05
    ----- 1.8/13.3 MB 4.2 MB/s eta
0:00:03
    ----- 2.9/13.3 MB 4.8 MB/s eta
```

```
0:00:03          ----- 3.9/13.3 MB 4.5 MB/s eta
0:00:03          ----- 5.2/13.3 MB 4.7 MB/s eta
0:00:02          ----- 6.6/13.3 MB 4.9 MB/s eta
0:00:02          ----- 7.9/13.3 MB 4.8 MB/s eta
0:00:02          ----- 9.2/13.3 MB 4.9 MB/s eta
0:00:01          ----- 10.0/13.3 MB 4.8 MB/s eta
0:00:01          ----- 10.5/13.3 MB 4.6 MB/s eta
0:00:01          ----- 11.0/13.3 MB 4.5 MB/s eta
0:00:01          ----- 12.1/13.3 MB 4.4 MB/s eta
0:00:01          ----- 13.1/13.3 MB 4.5 MB/s eta
0:00:01          ----- 13.3/13.3 MB 4.2 MB/s eta
0:00:00
Downloading cmdstanpy-1.2.5-py3-none-any.whl (94 kB)
Downloading holidays-0.70-py3-none-any.whl (903 kB)
----- 0.0/903.1 kB ? eta -:-:-
----- 903.1/903.1 kB 4.5 MB/s
eta 0:00:00
Downloading importlib_resources-6.5.2-py3-none-any.whl (37 kB)
Downloading stanio-0.5.1-py3-none-any.whl (8.1 kB)
Installing collected packages: stanio, importlib-resources, holidays,
cmdstanpy, Prophet
Successfully installed Prophet-1.1.6 cmdstanpy-1.2.5 holidays-0.70
importlib-resources-6.5.2 stanio-0.5.1

from prophet import Prophet

df = df.rename(columns={'Date': 'ds', 'Close': 'y'})
df.dropna(inplace=True)

train_size = int(len(df) * 0.8)
train = df.iloc[:train_size]
test = df.iloc[train_size:]

model = Prophet()
model.fit(train)

01:05:05 - cmdstanpy - INFO - Chain [1] start processing
01:05:10 - cmdstanpy - INFO - Chain [1] done processing

<prophet.forecaster.Prophet at 0x222675295e0>
```

```
future = model.make_future_dataframe(periods=len(test))
```