

Assignment 5

Name: Khalate Shubham

Class: AIDS A

Rollno: 23107062

Batch: C

Linear Regression

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

/home/admin1/anaconda3/lib/python3.9/site-packages/scipy/
__init__.py:132: UserWarning: A NumPy version >=1.21.6 and <1.28.0 is
required for this version of SciPy (detected version 1.21.5)
    warnings.warn(f"A NumPy version >={np_minversion} and
<{np_maxversion}")


df=pd.read_csv("BostonHousing.csv")
df
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222
..
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273

```

      ptratio      b    lstat   medv
0       15.3  396.90    4.98  24.0
1       17.8  396.90    9.14  21.6
2       17.8  392.83    4.03  34.7
3       18.7  394.63    2.94  33.4
4       18.7  396.90    5.33  36.2
..     ...
501     21.0  391.99    9.67  22.4
502     21.0  396.90    9.08  20.6
503     21.0  396.90    5.64  23.9
504     21.0  393.45    6.48  22.0
505     21.0  396.90    7.88  11.9

[506 rows x 14 columns]

df.shape
(506, 14)

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column   Non-Null Count  Dtype  
---  -- 
 0   crim     506 non-null    float64
 1   zn        506 non-null    float64
 2   indus    506 non-null    float64
 3   chas     506 non-null    int64  
 4   nox      506 non-null    float64
 5   rm        501 non-null    float64
 6   age       506 non-null    float64
 7   dis       506 non-null    float64
 8   rad       506 non-null    int64  
 9   tax       506 non-null    int64  
 10  ptratio   506 non-null    float64
 11  b         506 non-null    float64
 12  lstat    506 non-null    float64
 13  medv     506 non-null    float64
dtypes: float64(11), int64(3)
memory usage: 55.5 KB

df.isnull()

      crim      zn  indus    chas    nox      rm     age     dis     rad
tax \
0   False  False  False  False  False  False  False  False  False
False
1   False  False  False  False  False  False  False  False  False
False

```

```
2    False  
False  
3    False  
False  
4    False  
False  
...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...  
...  
501   False  
False  
502   False  
False  
503   False  
False  
504   False  
False  
505   False  
False
```

```
      ptratio      b      lstat      medv  
0    False  False  False  False  
1    False  False  False  False  
2    False  False  False  False  
3    False  False  False  False  
4    False  False  False  False  
...  ...  ...  ...  
501   False  False  False  False  
502   False  False  False  False  
503   False  False  False  False  
504   False  False  False  False  
505   False  False  False  False
```

[506 rows x 14 columns]

df.describe()

	crim	zn	indus	chas	nox
rm \ count	506.000000	506.000000	506.000000	506.000000	506.000000
501.000000					
mean	3.613524	11.363636	11.136779	0.069170	0.554695
6.284341					
std	8.601545	23.322453	6.860353	0.253994	0.115878
0.705587					
min	0.006320	0.000000	0.460000	0.000000	0.385000
3.561000					
25%	0.082045	0.000000	5.190000	0.000000	0.449000
5.884000					
50%	0.256510	0.000000	9.690000	0.000000	0.538000
6.208000					

```
75%      3.677083   12.500000   18.100000   0.000000   0.624000  
6.625000  
max     88.976200 100.000000  27.740000   1.000000   0.871000  
8.780000
```

```
          age       dis       rad       tax      ptratio  
b \  
count 506.000000 506.000000 506.000000 506.000000 506.000000  
506.000000  
mean  68.574901  3.795043  9.549407 408.237154 18.455534  
356.674032  
std   28.148861  2.105710  8.707259 168.537116 2.164946  
91.294864  
min   2.900000  1.129600  1.000000 187.000000 12.600000  
0.320000  
25%   45.025000  2.100175  4.000000 279.000000 17.400000  
375.377500  
50%   77.500000  3.207450  5.000000 330.000000 19.050000  
391.440000  
75%   94.075000  5.188425  24.000000 666.000000 20.200000  
396.225000  
max   100.000000 12.126500 24.000000 711.000000 22.000000  
396.900000
```

```
          lstat      medv  
count 506.000000 506.000000  
mean 12.653063 22.532806  
std  7.141062 9.197104  
min  1.730000 5.000000  
25%  6.950000 17.025000  
50% 11.360000 21.200000  
75% 16.955000 25.000000  
max 37.970000 50.000000
```

```
df.head()
```

```
      crim      zn  indus  chas      nox      rm      age      dis      rad      tax  
ptratio \  
0 0.00632  18.0  2.31      0  0.538  6.575  65.2  4.0900  1  296  
15.3  
1 0.02731  0.0  7.07      0  0.469  6.421  78.9  4.9671  2  242  
17.8  
2 0.02729  0.0  7.07      0  0.469  7.185  61.1  4.9671  2  242  
17.8  
3 0.03237  0.0  2.18      0  0.458  6.998  45.8  6.0622  3  222  
18.7  
4 0.06905  0.0  2.18      0  0.458  7.147  54.2  6.0622  3  222  
18.7
```

```
      b      lstat      medv
```

```
0 396.90 4.98 24.0
1 396.90 9.14 21.6
2 392.83 4.03 34.7
3 394.63 2.94 33.4
4 396.90 5.33 36.2
```

```
df.dtypes
```

```
crim      float64
zn        float64
indus     float64
chas      int64
nox       float64
rm        float64
age       float64
dis       float64
rad       int64
tax       int64
ptratio   float64
b         float64
lstat     float64
medv     float64
dtype: object
```

```
df.tail()
```

```
      crim    zn  indus  chas    nox    rm    age    dis    rad    tax
ptratio \
501  0.06263  0.0  11.93     0  0.573  6.593  69.1  2.4786    1  273
21.0
502  0.04527  0.0  11.93     0  0.573  6.120  76.7  2.2875    1  273
21.0
503  0.06076  0.0  11.93     0  0.573  6.976  91.0  2.1675    1  273
21.0
504  0.10959  0.0  11.93     0  0.573  6.794  89.3  2.3889    1  273
21.0
505  0.04741  0.0  11.93     0  0.573  6.030  80.8  2.5050    1  273
21.0
```

```
      b    lstat  medv
501 391.99  9.67 22.4
502 396.90  9.08 20.6
503 396.90  5.64 23.9
504 393.45  6.48 22.0
505 396.90  7.88 11.9
```

```
df.notnull()
```

```
      crim    zn  indus  chas    nox    rm    age    dis    rad    tax
ptratio \
0    True   True   True   True   True   True   True   True   True   True
```

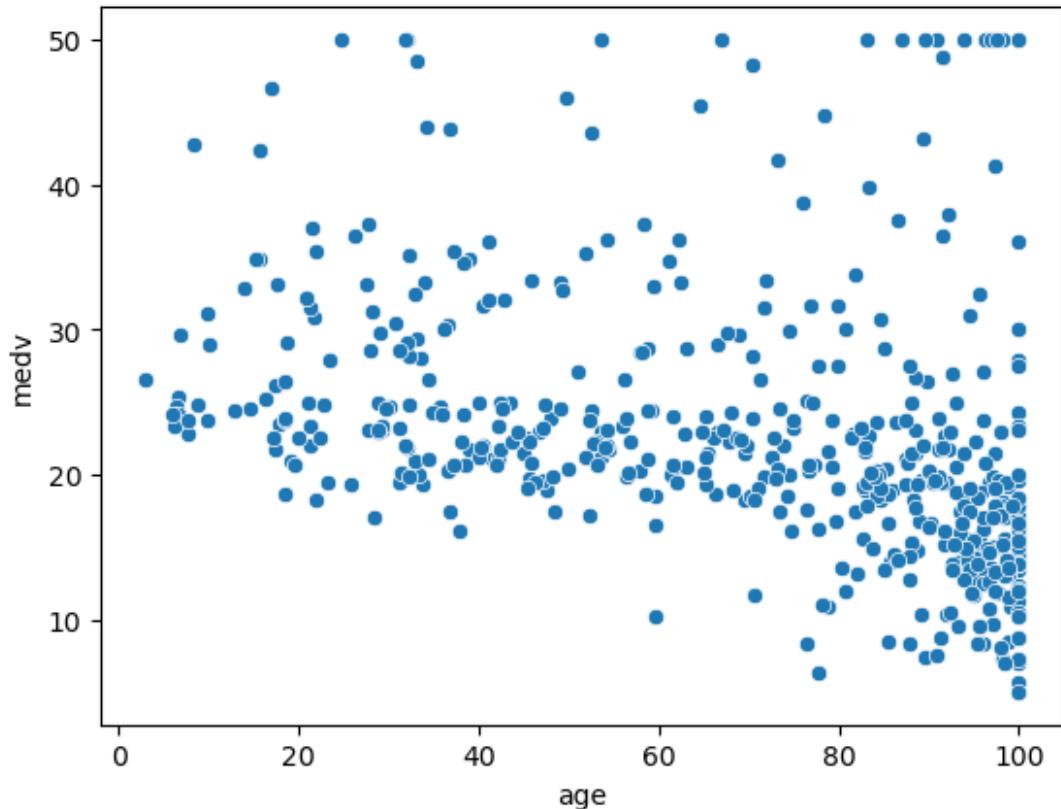
```
True
1  True  True   True  True  True  True  True  True  True  True  True
True
2  True  True   True  True  True  True  True  True  True  True  True
True
3  True  True   True  True  True  True  True  True  True  True  True
True
4  True  True   True  True  True  True  True  True  True  True  True
True
...
...
501 True  True   True  True  True  True  True  True  True  True  True
True
502 True  True   True  True  True  True  True  True  True  True  True
True
503 True  True   True  True  True  True  True  True  True  True  True
True
504 True  True   True  True  True  True  True  True  True  True  True
True
505 True  True   True  True  True  True  True  True  True  True  True
True

      b  lstat  medv
0  True  True  True
1  True  True  True
2  True  True  True
3  True  True  True
4  True  True  True
...
...
501 True  True  True
502 True  True  True
503 True  True  True
504 True  True  True
505 True  True  True

[506 rows x 14 columns]

sns.scatterplot(x=df['age'],y=df['medv'])

<Axes: xlabel='age', ylabel='medv'>
```



```
x=df['medv'].drop  
y=df['medv']  
  
x  
  
<bound method Series.drop of 0      24.0  
1      21.6  
2      34.7  
3      33.4  
4      36.2  
...  
501     22.4  
502     20.6  
503     23.9  
504     22.0  
505     11.9  
Name: medv, Length: 506, dtype: float64>  
  
y  
  
0      24.0  
1      21.6  
2      34.7  
3      33.4  
4      36.2
```

```
    ...
501    22.4
502    20.6
503    23.9
504    22.0
505    11.9
Name: medv, Length: 506, dtype: float64

from sklearn import linear_model
lr=linear_model.LinearRegression()
lr
LinearRegression()

x=df.drop(columns=[ 'medv' ])
y=df[ 'medv' ]

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y,
train_size=0.8, test_size=0.2, random_state=5)

len(x_train)
404

len(y_test)
102

X_cleaned = x.dropna()
y_cleaned = y[X_cleaned.index]

lr.fit(x_test,y_test)
LinearRegression()

y_predict=lr.predict(x_test)
y_predict
array([38.6519417 , 29.30271169, 26.47847516,  8.03513159,
37.42476358,
       10.78384061, 27.22276743, 30.21226213, 26.61253562,
21.06652464,
       34.02779035, 20.21309397, 21.71723439, 33.40530263,
27.62503723,
       16.47311571,  4.98046603, 16.03843732, 14.40028301,
18.32973666,
       7.2078973 , 19.95611372, 40.62816672, 24.36936464,
```

```
33.29161476,
           13.46962143, 23.79116653, 22.67376479, 21.5233983 ,
21.23433839,
           20.2741073 , 10.42265123, 16.56437975, 25.66102109,
28.75403451,
           19.66675985, 28.81834008, 13.36087535, 44.14506966,
32.54607535,
           18.5855062 , 8.45539158, 28.06191383, 12.75079081,
28.17291196,
           30.48034899, 4.63763795, 21.93115392, 21.31022511,
17.76197326,
           20.14192739, 20.33332662, 23.87349206, 16.79357584,
17.43706481,
           23.77962221, 37.77957524, 17.22313246, 29.70530794,
22.13612898,
           19.92467917, 24.32205673, 14.97385835, 32.11376975,
18.1217244 ,
           13.19370126, 20.76593993, 25.11864473, 21.92814514,
17.86530713,
           20.64138587, 24.31057453, 18.11694057, 23.38517598,
15.47640546,
           25.4734131 , 22.33367259, 16.75472738, 36.06092324,
16.03235008,
           20.00420155, 44.21358346, 24.84747351, 16.34231002,
23.02251201,
           18.15512877, 16.78925331, 9.07444408, 21.23859601,
16.5220274 ,
           38.18286199, 17.4568254 , 20.5909679 , 17.26019255,
24.63250975,
           27.59841206, 14.00680196, 24.37225018, 22.46185284,
13.1520875 ,
           20.74853494, 22.00255808])
```

```
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x_test, y_test)

LinearRegression()

lr.fit(x_test,y_test)

LinearRegression()

x_predict=lr.predict(x_test)

test_score = model.score(x_test, y_test)

print("Accuracy score:", test_score)

Accuracy score: 0.8177826690284403
```

```
from sklearn.metrics import
mean_squared_error,mean_absolute_error,r2_score
mse=mean_squared_error(y_test,x_predict)
mse
14.266499786743381

r2=r2_score(y_test,x_predict)
r2
0.8177826690284403

mae=mean_absolute_error(y_test,x_predict)
mae
2.8858058614979654

import matplotlib.pyplot as plt

plt.scatter(y_test,x_predict,color='red',edgecolor='k',alpha=0.7)
plt.plot([min(y_test),max(y_test)],
[min(y_test),max(y_test)],color='purple',linestyle="--")
plt.xlabel("actual values((medv))")
plt.ylabel("predicted values((medv))")
plt.title("actual vs predicted vaues with best fit line")
plt.show()
```

actual vs predicted values with best fit line

